

# landmark

January 20, 2022

## 1 Convolutional Neural Networks

### 1.1 Project: Write an Algorithm for Landmark Classification

---

In this notebook, some template code has already been provided for you, and you will need to implement additional functionality to successfully complete this project. You will not need to modify the included code beyond what is requested. Sections that begin with '**(IMPLEMENTATION)**' in the header indicate that the following block of code will require additional functionality which you must provide. Instructions will be provided for each section, and the specifics of the implementation are marked in the code block with a 'TODO' statement. Please be sure to read the instructions carefully!

**Note:** Once you have completed all the code implementations, you need to finalize your work by exporting the Jupyter Notebook as an HTML document. Before exporting the notebook to HTML, all the code cells need to have been run so that reviewers can see the final implementation and output. You can then export the notebook by using the menu above and navigating to **File -> Download as -> HTML (.html)**. Include the finished document along with this notebook as your submission.

In addition to implementing code, there will be questions that you must answer which relate to the project and your implementation. Each section where you will answer a question is preceded by a '**Question X**' header. Carefully read each question and provide thorough answers in the following text boxes that begin with '**Answer:**'. Your project submission will be evaluated based on your answers to each of the questions and the implementation you provide.

**Note:** Code and Markdown cells can be executed using the **Shift + Enter** keyboard shortcut. Markdown cells can be edited by double-clicking the cell to enter edit mode.

The rubric contains *optional* "Stand Out Suggestions" for enhancing the project beyond the minimum requirements. If you decide to pursue the "Stand Out Suggestions", you should include the code in this Jupyter notebook.

## Step 0: Download Datasets and Install Python Modules

**Note:** if you are using the Udacity workspace, *YOU CAN SKIP THIS STEP*. The dataset can be found in the /data folder and all required Python modules have been installed in the workspace.

Download the [landmark dataset](#). Unzip the folder and place it in this project's home directory, at the location /landmark\_images.

Install the following Python modules: \* cv2 \* matplotlib \* numpy \* PIL \* torch \* torchvision

---

### ## Step 1: Create a CNN to Classify Landmarks (from Scratch)

In this step, you will create a CNN that classifies landmarks. You must create your CNN *from scratch* (so, you can't use transfer learning *yet!*), and you must attain a test accuracy of at least 20%.

Although 20% may seem low at first glance, it seems more reasonable after realizing how difficult of a problem this is. Many times, an image that is taken at a landmark captures a fairly mundane image of an animal or plant, like in the following picture.

Just by looking at that image alone, would you have been able to guess that it was taken at the Haleakal National Park in Hawaii?

An accuracy of 20% is significantly better than random guessing, which would provide an accuracy of just 2%. In Step 2 of this notebook, you will have the opportunity to greatly improve accuracy by using transfer learning to create a CNN.

Remember that practice is far ahead of theory in deep learning. Experiment with many different architectures, and trust your intuition. And, of course, have fun!

#### 1.1.1 (IMPLEMENTATION) Specify Data Loaders for the Landmark Dataset

Use the code cell below to create three separate [data loaders](#): one for training data, one for validation data, and one for test data. Randomly split the images located at `landmark_images/train` to create the train and validation data loaders, and use the images located at `landmark_images/test` to create the test data loader.

**Note:** Remember that the dataset can be found at `/data/landmark_images/` in the workspace.

All three of your data loaders should be accessible via a dictionary named `loaders_scratch`. Your train data loader should be at `loaders_scratch['train']`, your validation data loader should be at `loaders_scratch['valid']`, and your test data loader should be at `loaders_scratch['test']`.

You may find [this documentation on custom datasets](#) to be a useful resource. If you are interested in augmenting your training and/or validation data, check out the wide variety of [transforms](#)!

```
In [1]: ### TODO: Write data loaders for training, validation, and test sets
        ## Specify appropriate transforms, and batch_sizes
        import torch
        import numpy as np

        import torchvision
        from torchvision import datasets, transforms
        from torch.utils.data.sampler import SubsetRandomSampler

        # Defining the data path
        data_path = '/data/landmark_images/'
        # Defining batch size to load
        batch_size = 32
        # % of train set to use in validation
        valid_percent = 0.2

        # Specifying the transforms
```

```

train_transforms = transforms.Compose([transforms.RandomRotation(30),
                                       transforms.RandomResizedCrop(224),
                                       transforms.RandomHorizontalFlip(),
                                       transforms.ToTensor(),
                                       #transforms.Normalize(0.5836, 0.1628) # zip argum
                                       transforms.Normalize((0.485, 0.456, 0.406), # Was
                                                           (0.229, 0.224, 0.225)) # Sol
                                       ])

test_transforms = transforms.Compose([transforms.Resize(255),
                                       transforms.CenterCrop(224),
                                       transforms.ToTensor(),
                                       #transforms.Normalize(0.5836, 0.1628)
                                       transforms.Normalize((0.485, 0.456, 0.406),
                                                           (0.229, 0.224, 0.225))
                                       ])

# Loading the datasets
train_data = datasets.ImageFolder(data_path + '/train', transform=train_transforms)
test_data = datasets.ImageFolder(data_path + '/test', transform=test_transforms)
print(len(train_data))
print(len(test_data))

# Splitting data to train and validation
num_train = len(train_data)
indices = list(range(num_train))
np.random.shuffle(indices)
split = int(np.floor(valid_percent * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# Defining samplers
train_sampler = SubsetRandomSampler(train_idx)
valid_sampler = SubsetRandomSampler(valid_idx)

# Preparing the data loaders
train_load_scr = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                              sampler=train_sampler)
valid_load_scr = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                              sampler=valid_sampler)
test_load_scr = torch.utils.data.DataLoader(test_data, batch_size=batch_size)

# Compressing in a dict
loaders_scratch = {'train': train_load_scr, 'valid': valid_load_scr, 'test': test_load_scr}

```

4996  
1250

**Question 1:** Describe your chosen procedure for preprocessing the data. - How does your code

resize the images (by cropping, stretching, etc)? What size did you pick for the input tensor, and why? - Did you decide to augment the dataset? If so, how (through translations, flips, rotations, etc)? If not, why not?

**Answer:** - I chose this 224 size because I read in discussions online that was a good resolution to train it without using too much memory and it was also the size used for the VGG transferred NN, so some time saved there just copying the code.

- My code resize the train data to 224x224 with a random resized crop included. It also rotate the data in a random direction, apart from flipping it horizontally. Then, the 3 rgb channels get normalized and the data converted to tensor.
- Yes, with this augmentations mentioned above in the training data, I tried to get higher accuracy both on the scratch NN and the transferred one. As suggested in the *Stand Out Suggestions 1*
- For the test set I just used a stretching resizer and the a center crop to 224 with no augmentations.

### 1.1.2 (IMPLEMENTATION) Visualize a Batch of Training Data

Use the code cell below to retrieve a batch of images from your train data loader, display at least 5 images simultaneously, and label each displayed image with its class name (e.g., "Golden Gate Bridge").

Visualizing the output of your data loader is a great way to ensure that your data loading and preprocessing are working as expected.

```
In [2]: import matplotlib.pyplot as plt
        %matplotlib inline

        ## TODO: visualize a batch of the train data loader
        # Obtaining one batch of training images
        dataiter = iter(loaders_scratch['train'])
        images, labels = dataiter.next()
        # print(images[0].mean(), images[0].std())
        print(images.size())
        images = images.numpy() # convert images to numpy for display
        ## the class names can be accessed at the `classes` attribute
        ## of your dataset object (e.g., `train_dataset.classes`)
        classes = train_data.classes
        # print(len(classes))
        # Plotting the images in the batch, along with the corresponding labels
        fig = plt.figure(figsize=(25, 4))
        for idx in np.arange(5):
            ax = fig.add_subplot(2, 8, idx+1, xticks=[], yticks=[])
            image = images[idx] / 2 + 0.5 # @Tejas J forum tip
            plt.imshow(np.transpose(image, (1, 2, 0)).clip(0,1)) # okay clip
            ax.set_title(classes[labels[idx]])

torch.Size([32, 3, 224, 224])
```

44.Trevi\_Fountain



03.Dead\_Sea




02.Ljubljana\_Castle



49.Temple\_of\_Olympian\_Zeus



45.Temple\_of\_Heaven



### 1.1.3 Initialize use\_cuda variable

```
In [3]: # useful variable that tells us whether we should use the GPU
        use_cuda = torch.cuda.is_available()
```

### 1.1.4 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_scratch`, and fill in the function `get_optimizer_scratch` below.

```
In [4]: ## TODO: select loss function
        import torch.nn as nn
        import torch.optim as optim

        criterion_scratch = nn.CrossEntropyLoss()

        def get_optimizer_scratch(model):
            ## TODO: select and return an optimizer
            optimizer = optim.Adamax(model.parameters(), lr=0.003)
            return optimizer
```

### 1.1.5 (IMPLEMENTATION) Model Architecture

Create a CNN to classify images of landmarks. Use the template in the code cell below.

```
In [5]: import torch.nn.functional as F
        # define the CNN architecture
        class Net(nn.Module):
            ## TODO: choose an architecture, and complete the class
            def __init__(self):
                super(Net, self).__init__()
                ## Define layers of a CNN
                self.conv1 = nn.Conv2d(3, 16, 3, padding=1)
                self.conv2 = nn.Conv2d(16, 32, 3, padding=1)
                self.conv3 = nn.Conv2d(32, 64, 3, padding=1)
                self.pool = nn.MaxPool2d(2, 2)

                self.fc1 = nn.Linear(64 * 28 * 28, 500)
                self.fc2 = nn.Linear(500, 50)
```

```

        self.dropout = nn.Dropout(0.20)

    def forward(self, x):
        ## Define forward behavior
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        # print(x.shape)
        x = x.view(-1, 64 * 28 * 28)

        x = self.dropout(x)
        x = F.relu(self.fc1(x))
        x = self.dropout(x)
        x = self.fc2(x)

    return x

### Do NOT modify the code below this line. ###

# instantiate the CNN
model_scratch = Net()
print(model_scratch)
# move tensors to GPU if CUDA is available
if use_cuda:
    model_scratch.cuda()

Net(
  (conv1): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (conv3): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (fc1): Linear(in_features=50176, out_features=500, bias=True)
  (fc2): Linear(in_features=500, out_features=50, bias=True)
  (dropout): Dropout(p=0.2)
)

```

**Question 2:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step.

**Answer:**

I didn't want to use a really big NN so it was going to need too much GPU time for train. So I chose a network similar to which I have used in my CIFAR exercise due to it was a similar multi-label task, in which I reached a good accuracy. Then I arranged the shapes of the data used so it will work good using the `print(x.shape)`.

Everything worked fine so quickly and I was training it to 40 epochs in a moment. Then the val loss stop decreasing so I stopped the training with a 44% accu.

### 1.1.6 (IMPLEMENTATION) Implement the Training Algorithm

Implement your training algorithm in the code cell below. Save the final model parameters at the filepath stored in the variable `save_path`.

```
In [6]: from tqdm import tqdm_notebook as tqdm
def train(n_epochs, loaders, model, optimizer, criterion, use_cuda, save_path):
    """returns trained model"""
    # initialize tracker for minimum validation loss
    valid_loss_min = np.Inf # Need a change for updating this whenever u want to continu

    for epoch in range(1, n_epochs+1):
        # initialize variables to monitor training and validation loss
        train_loss = 0.0
        valid_loss = 0.0

        #####
        # train the model #
        #####
        # set the module to training mode
        model.train()
        for batch_idx, (data, target) in tqdm(enumerate(loaders['train'])):
            # move to GPU
            if use_cuda:
                data, target = data.cuda(), target.cuda()

            ## TODO: find the loss and update the model parameters accordingly
            ## record the average training loss, using something like
            ## train_loss = train_loss + ((1 / (batch_idx + 1)) * (loss.data.item() - tr

            # First, reset the optimizer
            optimizer.zero_grad()
            # Second, pass the data through the model
            output = model(data)
            # Calculate the loss
            loss = criterion(output, target)
            # Backpropagate the losses
            loss.backward()
            # Optimize the parameters
            optimizer.step()
            # Update the train loss
            train_loss = train_loss + (1 / (batch_idx + 1)) * (loss.data.item() - train_

        #####
        # validate the model #
        #####
        # set the model to evaluation mode
        model.eval()
        for batch_idx, (data, target) in tqdm(enumerate(loaders['valid'])):
```

```

        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()

        ## TODO: update average validation loss

        # Pass the input to the model
        output = model(data)
        # Calculate the loss
        loss = criterion(output, target)
        # Update the validation loss
        valid_loss = valid_loss + (1 / (batch_idx + 1)) * (loss.data.item() - valid_loss)

    # print training/validation statistics
    print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
        epoch,
        train_loss,
        valid_loss
    ))

    ## TODO: if the validation loss has decreased, save the model at the filepath specified
    if valid_loss <= valid_loss_min:
        print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
            valid_loss_min,
            valid_loss))
        torch.save(model.state_dict(), save_path)
        valid_loss_min = valid_loss

return model

```

### 1.1.7 (IMPLEMENTATION) Experiment with the Weight Initialization

Use the code cell below to define a custom weight initialization, and then train with your weight initialization for a few epochs. Make sure that neither the training loss nor validation loss is nan.

Later on, you will be able to see how this compares to training with PyTorch's default weight initialization.

```

In [19]: def custom_weight_init(m):
        ## TODO: implement a weight initialization strategy
        classname = m.__class__.__name__
        if classname.find('Linear') != -1:
            n = m.in_features
            y = (1.01234/np.sqrt(n))
            m.weight.data.normal_(0.0012321, y)
            m.bias.data.fill_(0.04129)
        elif classname.find('Conv2d') != -1:

```



```

        torch.nn.init.kaiming_normal_(m.weight)
        if m.bias is not None:
            m.bias.data.fill_(0.014184)

### Do NOT modify the code below this line. ###

model_scratch.apply(custom_weight_init)
model_scratch = train(20, loaders_scratch, model_scratch, get_optimizer_scratch(model_s
                    criterion_scratch, use_cuda, 'ignore.pt')

HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))


HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))


Epoch: 1          Training Loss: 9.263409          Validation Loss: 3.921264
Validation loss decreased (inf --> 3.921264).  Saving model ...

HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))


HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))


Epoch: 2          Training Loss: 3.911889          Validation Loss: 3.915108
Validation loss decreased (3.921264 --> 3.915108).  Saving model ...

HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))


HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))

```

Epoch: 3            Training Loss: 3.908436            Validation Loss: 3.912008  
Validation loss decreased (3.915108 --> 3.912008). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 4            Training Loss: 3.904149            Validation Loss: 3.894633  
Validation loss decreased (3.912008 --> 3.894633). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 5            Training Loss: 3.884599            Validation Loss: 3.875511  
Validation loss decreased (3.894633 --> 3.875511). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 6            Training Loss: 3.856737            Validation Loss: 3.845820  
Validation loss decreased (3.875511 --> 3.845820). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 7          Training Loss: 3.827743          Validation Loss: 3.816581  
Validation loss decreased (3.845820 --> 3.816581). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 8          Training Loss: 3.780379          Validation Loss: 3.752029  
Validation loss decreased (3.816581 --> 3.752029). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 9          Training Loss: 3.758815          Validation Loss: 3.744243  
Validation loss decreased (3.752029 --> 3.744243). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 10         Training Loss: 3.730066          Validation Loss: 3.701447  
Validation loss decreased (3.744243 --> 3.701447). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

-----

KeyboardInterrupt

Traceback (most recent call last)

```
<ipython-input-19-ca9ac84c34a2> in <module>()
    17 model_scratch.apply(custom_weight_init)
    18 model_scratch = train(20, loaders_scratch, model_scratch, get_optimizer_scratch(model_scratch),
--> 19         criterion_scratch, use_cuda, 'ignore.pt')

<ipython-input-6-9fcd06033abe> in train(n_epochs, loaders, model, optimizer, criterion,
    15     # set the module to training mode
    16     model.train()
--> 17     for batch_idx, (data, target) in tqdm(enumerate(loaders['train'])):
    18         # move to GPU
    19         if use_cuda:

/opt/conda/lib/python3.6/site-packages/tqdm/_tqdm_notebook.py in __iter__(self, *args, **kwargs)
    185     def __iter__(self, *args, **kwargs):
    186         try:
--> 187             for obj in super(tqdm_notebook, self).__iter__(*args, **kwargs):
    188                 # return super(tqdm...) will not catch exception
    189                 yield obj

/opt/conda/lib/python3.6/site-packages/tqdm/_tqdm.py in __iter__(self)
    831     """", fp_write=getattr(self.fp, 'write', sys.stderr.write))
    832
--> 833     for obj in iterable:
    834         yield obj
    835         # Update and print the progressbar.

/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py in __next__(self)
    262     if self.num_workers == 0: # same-process loading
    263         indices = next(self.sample_iter) # may raise StopIteration
--> 264         batch = self.collate_fn([self.dataset[i] for i in indices])
    265         if self.pin_memory:
    266             batch = pin_memory_batch(batch)

/opt/conda/lib/python3.6/site-packages/torch/utils/data/dataloader.py in <listcomp>(.0)
    262     if self.num_workers == 0: # same-process loading
```

```

263             indices = next(self.sample_iter) # may raise StopIteration
--> 264             batch = self.collate_fn([self.dataset[i] for i in indices])
265             if self.pin_memory:
266                 batch = pin_memory_batch(batch)

/opt/conda/lib/python3.6/site-packages/torchvision-0.2.1-py3.6.egg/torchvision/datasets/
99         """
100         path, target = self.samples[index]
--> 101         sample = self.loader(path)
102         if self.transform is not None:
103             sample = self.transform(sample)

/opt/conda/lib/python3.6/site-packages/torchvision-0.2.1-py3.6.egg/torchvision/datasets/
145         return accimage_loader(path)
146     else:
--> 147         return pil_loader(path)
148
149

/opt/conda/lib/python3.6/site-packages/torchvision-0.2.1-py3.6.egg/torchvision/datasets/
128     with open(path, 'rb') as f:
129         img = Image.open(f)
--> 130         return img.convert('RGB')
131
132

/opt/conda/lib/python3.6/site-packages/PIL/Image.py in convert(self, mode, matrix, dithering)
890         """
891
--> 892         self.load()
893
894         if not mode and self.mode == "P":

/opt/conda/lib/python3.6/site-packages/PIL/ImageFile.py in load(self)
233
234         b = b + s
--> 235         n, err_code = decoder.decode(b)
236         if n < 0:
237             break

```

KeyboardInterrupt:

### 1.1.8 (IMPLEMENTATION) Train and Validate the Model

Run the next code cell to train your model.

```
In [7]: ## TODO: you may change the number of epochs if you'd like,  
        ## but changing it is not required  
        num_epochs = 100  
  
        ##-## Do NOT modify the code below this line. ##-##  
  
        # function to re-initialize a model with pytorch's default weight initialization  
        def default_weight_init(m):  
            reset_parameters = getattr(m, 'reset_parameters', None)  
            if callable(reset_parameters):  
                m.reset_parameters()  
  
        # reset the model parameters  
        model_scratch.apply(default_weight_init)  
  
        # train the model  
        model_scratch.load_state_dict(torch.load('model_scratch.pt'))  
        model_scratch = train(num_epochs, loaders_scratch, model_scratch, get_optimizer_scratch(  
                                criterion_scratch, use_cuda, 'model_scratch.pt'))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 1           Training Loss: 4.075562           Validation Loss: 3.888436  
Validation loss decreased (inf --> 3.888436). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 2           Training Loss: 3.839315           Validation Loss: 3.774903  
Validation loss decreased (3.888436 --> 3.774903). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 3          Training Loss: 3.747534          Validation Loss: 3.737884  
Validation loss decreased (3.774903 --> 3.737884). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 4          Training Loss: 3.691145          Validation Loss: 3.653192  
Validation loss decreased (3.737884 --> 3.653192). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 5          Training Loss: 3.628869          Validation Loss: 3.545792  
Validation loss decreased (3.653192 --> 3.545792). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 6            Training Loss: 3.561285            Validation Loss: 3.534850  
Validation loss decreased (3.545792 --> 3.534850). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 7            Training Loss: 3.464760            Validation Loss: 3.374228  
Validation loss decreased (3.534850 --> 3.374228). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 8            Training Loss: 3.406613            Validation Loss: 3.353082  
Validation loss decreased (3.374228 --> 3.353082). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 9            Training Loss: 3.346248            Validation Loss: 3.275855  
Validation loss decreased (3.353082 --> 3.275855). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))



```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 10          Training Loss: 3.281782          Validation Loss: 3.225167  
Validation loss decreased (3.275855 --> 3.225167). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 11          Training Loss: 3.188290          Validation Loss: 3.200020  
Validation loss decreased (3.225167 --> 3.200020). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 12          Training Loss: 3.145452          Validation Loss: 3.114134  
Validation loss decreased (3.200020 --> 3.114134). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 13          Training Loss: 3.070233          Validation Loss: 3.113539  
Validation loss decreased (3.114134 --> 3.113539). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 14          Training Loss: 3.004590          Validation Loss: 3.025699  
Validation loss decreased (3.113539 --> 3.025699). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 15          Training Loss: 2.983421          Validation Loss: 3.092224
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 16          Training Loss: 2.903275          Validation Loss: 2.942105  
Validation loss decreased (3.025699 --> 2.942105). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 17            Training Loss: 2.827653            Validation Loss: 2.931641  
Validation loss decreased (2.942105 --> 2.931641). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 18            Training Loss: 2.804891            Validation Loss: 2.940915

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 19            Training Loss: 2.747423            Validation Loss: 2.882969  
Validation loss decreased (2.931641 --> 2.882969). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 20            Training Loss: 2.719472            Validation Loss: 2.811672  
Validation loss decreased (2.882969 --> 2.811672). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 21                    Training Loss: 2.662407                    Validation Loss: 2.844345

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 22                    Training Loss: 2.616170                    Validation Loss: 2.731724  
Validation loss decreased (2.811672 --> 2.731724). Saving model ...

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 23                    Training Loss: 2.588323                    Validation Loss: 2.787043

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 24                    Training Loss: 2.552168                    Validation Loss: 2.741770

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 25           Training Loss: 2.527035           Validation Loss: 2.790179
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 26           Training Loss: 2.465996           Validation Loss: 2.724202  
Validation loss decreased (2.731724 --> 2.724202). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 27           Training Loss: 2.500872           Validation Loss: 2.685225  
Validation loss decreased (2.724202 --> 2.685225). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 28           Training Loss: 2.446638           Validation Loss: 2.685790
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 29          Training Loss: 2.385067          Validation Loss: 2.704247
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 30          Training Loss: 2.369251          Validation Loss: 2.655194  
Validation loss decreased (2.685225 --> 2.655194). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 31          Training Loss: 2.336515          Validation Loss: 2.615879  
Validation loss decreased (2.655194 --> 2.615879). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

Epoch: 32            Training Loss: 2.288093            Validation Loss: 2.586198  
Validation loss decreased (2.615879 --> 2.586198). Saving model ...

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 33            Training Loss: 2.250560            Validation Loss: 2.629656

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 34            Training Loss: 2.212519            Validation Loss: 2.662029

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

Epoch: 35            Training Loss: 2.200109            Validation Loss: 2.676562

HBox(children=(IntProgress(value=1, bar\_style='info', max=1), HTML(value='')))

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 36           Training Loss: 2.184821           Validation Loss: 2.535587  
Validation loss decreased (2.586198 --> 2.535587). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 37           Training Loss: 2.123658           Validation Loss: 2.652949
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 38           Training Loss: 2.137784           Validation Loss: 2.654183
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 39           Training Loss: 2.077990           Validation Loss: 2.660160
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```



```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 40          Training Loss: 2.090369          Validation Loss: 2.511839
Validation loss decreased (2.535587 --> 2.511839).  Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 41          Training Loss: 2.049895          Validation Loss: 2.651264
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
-----
KeyboardInterrupt                                Traceback (most recent call last)

<ipython-input-7-234244f91f3a> in <module>()
    16 # train the model
    17 model_scratch = train(num_epochs, loaders_scratch, model_scratch, get_optimizer_scratch,
---> 18                          criterion_scratch, use_cuda, 'model_scratch.pt')

<ipython-input-6-9fcd06033abe> in train(n_epochs, loaders, model, optimizer, criterion,
    35         optimizer.step()
    36         # Update the train loss
---> 37         train_loss = train_loss + (1 / (batch_idx + 1)) * (loss.data.item() - tr
    38
    39         #####
```

```
KeyboardInterrupt:
```

### 1.1.9 (IMPLEMENTATION) Test the Model

Run the code cell below to try out your model on the test dataset of landmark images. Run the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 20%.

```
In [7]: def test(loaders, model, criterion, use_cuda):

    # monitor test loss and accuracy
    test_loss = 0.
    correct = 0.
    total = 0.

    # set the module to evaluation mode
    model.eval()

    for batch_idx, (data, target) in enumerate(loaders['test']):
        # move to GPU
        if use_cuda:
            data, target = data.cuda(), target.cuda()
        # forward pass: compute predicted outputs by passing inputs to the model
        output = model(data)
        # calculate the loss
        loss = criterion(output, target)
        # update average test loss
        test_loss = test_loss + ((1 / (batch_idx + 1)) * (loss.data.item() - test_loss))
        # convert output probabilities to predicted class
        pred = output.data.max(1, keepdim=True)[1]
        # compare predictions to true label
        correct += np.sum(np.squeeze(pred.eq(target.data.view_as(pred))).cpu().numpy())
        total += data.size(0)

    print('Test Loss: {:.6f}\n'.format(test_loss))

    print('\nTest Accuracy: %2d%% (%2d/%2d)' % (
        100. * correct / total, correct, total))

In [8]: # load the model that got the best validation accuracy
model_scratch.load_state_dict(torch.load('model_scratch.pt'))
test(loaders_scratch, model_scratch, criterion_scratch, use_cuda)
```

Test Loss: 2.305133

Test Accuracy: 44% (551/1250)

## Step 2: Create a CNN to Classify Landmarks (using Transfer Learning)

You will now use transfer learning to create a CNN that can identify landmarks from images. Your CNN must attain at least 60% accuracy on the test set.

### 1.1.10 (IMPLEMENTATION) Specify Data Loaders for the Landmark Dataset

Use the code cell below to create three separate [data loaders](#): one for training data, one for validation data, and one for test data. Randomly split the images located at `landmark_images/train` to create the train and validation data loaders, and use the images located at `landmark_images/test` to create the test data loader.

All three of your data loaders should be accessible via a dictionary named `loaders_transfer`. Your train data loader should be at `loaders_transfer['train']`, your validation data loader should be at `loaders_transfer['valid']`, and your test data loader should be at `loaders_transfer['test']`.

If you like, you are welcome to use the same data loaders from the previous step, when you created a CNN from scratch.

```
In [8]: ### TODO: Write data loaders for training, validation, and test sets
        ## Specify appropriate transforms, and batch_sizes
        # Defining the data path
        data_path = '/data/landmark_images/'
        # Defining batch size to load
        batch_size = 32
        # % of train set to use in validation
        valid_percent = 0.2

        # Specifying the transforms
        train_transforms = transforms.Compose([transforms.RandomRotation(30),
                                                transforms.RandomResizedCrop(224),
                                                transforms.RandomHorizontalFlip(),
                                                transforms.ToTensor(),
                                                #transforms.Normalize(0.5836, 0.1628) # zip argum
                                                transforms.Normalize((0.485, 0.456, 0.406), # Was
                                                                    (0.229, 0.224, 0.225)) # Sol
                                                ])

        test_transforms = transforms.Compose([transforms.Resize(255),
                                                transforms.CenterCrop(224),
                                                transforms.ToTensor(),
                                                #transforms.Normalize(0.5836, 0.1628)
                                                transforms.Normalize((0.485, 0.456, 0.406),
                                                                    (0.229, 0.224, 0.225))
                                                ])

        # Loading the datasets
        train_data = datasets.ImageFolder(data_path + '/train', transform=train_transforms)
        test_data = datasets.ImageFolder(data_path + '/test', transform=test_transforms)
        print(len(train_data))
```

```

print(len(test_data))

# Splitting data to train and validation
num_train = len(train_data)
indices = list(range(num_train))
np.random.shuffle(indices)
split = int(np.floor(valid_percent * num_train))
train_idx, valid_idx = indices[split:], indices[:split]

# Defining samplers
train_sampler = SubsetRandomSampler(train_idx)
valid_sampler = SubsetRandomSampler(valid_idx)

# Preparing the data loaders
train_load_tr = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             sampler=train_sampler)
valid_load_tr = torch.utils.data.DataLoader(train_data, batch_size=batch_size,
                                             sampler=valid_sampler)
test_load_tr = torch.utils.data.DataLoader(test_data, batch_size=batch_size)

loaders_transfer = {'train': train_load_tr, 'valid': valid_load_tr, 'test': test_load_tr}
4996
1250

```

### 1.1.11 (IMPLEMENTATION) Specify Loss Function and Optimizer

Use the next code cell to specify a [loss function](#) and [optimizer](#). Save the chosen loss function as `criterion_transfer`, and fill in the function `get_optimizer_transfer` below.

```

In [9]: ## TODO: select loss function
from torch import optim
criterion_transfer = nn.CrossEntropyLoss()

def get_optimizer_transfer(model):
    ## TODO: select and return optimizer
    optimizer = optim.Adamax(model.classifier.parameters(), lr=0.003)
    return optimizer

```

### 1.1.12 (IMPLEMENTATION) Model Architecture

Use transfer learning to create a CNN to classify images of landmarks. Use the code cell below, and save your initialized model as the variable `model_transfer`.

```

In [10]: ## TODO: Specify model architecture
import torch.nn as nn
from torchvision import models
from collections import OrderedDict

```

```

model_transfer = models.vgg19(pretrained=True)

for param in model_transfer.features.parameters():
    param.requires_grad_(False)

print(model_transfer)
#print(model_transfer.fc.in_features) # Inception and ResNet50 werent working properly
#print(model_transfer.fc.out_features)

#classifier = nn.Sequential(OrderedDict([ # VGG is more comfortable, I didnt even need
#                                     ('fc1', nn.Linear(model_transfer.fc.in_features, 1000)),
#                                     ('relu', nn.ReLU()),
#                                     ('drop', nn.Dropout(p=0.2)),
#                                     ('fc2', nn.Linear(1000, 500)),
#                                     ('relu2', nn.ReLU()),
#                                     ('drop2', nn.Dropout(p=0.2)),
#                                     ('fc3', nn.Linear(500, 50)),
#                                     ('output', nn.Softmax(dim=1))
#                                     ]))

model_transfer.classifier[6] = nn.Linear(in_features=4096, out_features=50)
#for param in model_transfer.fc.parameters(): # Code used with ResNet50 trying to unfreeze
#    param.requires_grad_(True)

print(model_transfer)

#-#-# Do NOT modify the code below this line. #-#-#

if use_cuda:
    model_transfer = model_transfer.cuda()

```

Downloading: "https://download.pytorch.org/models/vgg19-dcbb9e9d.pth" to /root/.torch/models/vgg19-dcbb9e9d.pth  
100%| 574673361/574673361 [00:06<00:00, 94410542.56it/s]

```

VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (8): ReLU(inplace)
    (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)

```

```

(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(17): ReLU(inplace)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=1000, bias=True)
)
)
VGG(
  (features): Sequential(
    (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU(inplace)
    (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU(inplace)
    (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (6): ReLU(inplace)
    (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))

```

```

(8): ReLU(inplace)
(9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(11): ReLU(inplace)
(12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(13): ReLU(inplace)
(14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(15): ReLU(inplace)
(16): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(17): ReLU(inplace)
(18): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(19): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(20): ReLU(inplace)
(21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(22): ReLU(inplace)
(23): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(24): ReLU(inplace)
(25): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(26): ReLU(inplace)
(27): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
(28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(29): ReLU(inplace)
(30): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(31): ReLU(inplace)
(32): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(33): ReLU(inplace)
(34): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
(35): ReLU(inplace)
(36): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
)
(classifier): Sequential(
  (0): Linear(in_features=25088, out_features=4096, bias=True)
  (1): ReLU(inplace)
  (2): Dropout(p=0.5)
  (3): Linear(in_features=4096, out_features=4096, bias=True)
  (4): ReLU(inplace)
  (5): Dropout(p=0.5)
  (6): Linear(in_features=4096, out_features=50, bias=True)
)
)

```

**Question 3:** Outline the steps you took to get to your final CNN architecture and your reasoning at each step. Describe why you think the architecture is suitable for the current problem.

**Answer:**

- I googled for the best available models in Torch 0.4.0 which is the version allowed in this workspace. And I tried the highest accu nets. First InceptionV3 was giving me a problem I wasn't managing to fix, so I decided to change to a ResNet50.

- With the ResNet I had 2 problems, first I was having problems with freezing some parameters but getting all frozen. Finally I managed to fix it but the net didn't learn well so I got a 30% accu what wasn't enough.
- Finally, I changed to the VGG19 which was easier to build, freeze and train. Changed the classifier so it fit my output classes and trained it for some epochs. It easily got pretty above the minimum accu.
- I think VGG is suitable because it was trained in a multilabel task with 1000 classes so the initial layers learned well how to extract the images features from a huge a varied dataset as ImageNet. I suppose it was easy to generalize to simple landmark images.

### 1.1.13 (IMPLEMENTATION) Train and Validate the Model

Train and validate your model in the code cell below. [Save the final model parameters](#) at filepath 'model\_transfer.pt'.

```
In [11]: # TODO: train the model and save the best model parameters at filepath 'model_transfer.pt'
num_epochs = 15

model_transfer.load_state_dict(torch.load('model_transfer.pt')) # For retraining after
model_transfer = train(num_epochs, loaders_transfer, model_transfer, get_optimizer_transfer,
                       criterion_transfer, use_cuda, 'model_transfer.pt')
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 1          Training Loss: 1.885878          Validation Loss: 1.233629
Validation loss decreased (inf --> 1.233629). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 2          Training Loss: 1.560930          Validation Loss: 1.258994
```



```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 3          Training Loss: 1.402346          Validation Loss: 1.154829  
Validation loss decreased (1.233629 --> 1.154829). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 4          Training Loss: 1.331803          Validation Loss: 1.103056  
Validation loss decreased (1.154829 --> 1.103056). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 5          Training Loss: 1.340903          Validation Loss: 1.093652  
Validation loss decreased (1.103056 --> 1.093652). Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 6           Training Loss: 1.291270           Validation Loss: 1.073554
Validation loss decreased (1.093652 --> 1.073554).  Saving model ...
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 7           Training Loss: 1.221157           Validation Loss: 1.084225
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

```
Epoch: 8           Training Loss: 1.190087           Validation Loss: 1.105835
```

```
HBox(children=(IntProgress(value=1, bar_style='info', max=1), HTML(value='')))
```

-----

```
KeyboardInterrupt                                Traceback (most recent call last)
```

```
<ipython-input-11-21b36c5e178c> in <module>()
      4 model_transfer.load_state_dict(torch.load('model_transfer.pt')) # For retraining aft
      5 model_transfer = train(num_epochs, loaders_transfer, model_transfer, get_optimizer_t
----> 6             criterion_transfer, use_cuda, 'model_transfer.pt')

<ipython-input-6-9fcd06033abe> in train(n_epochs, loaders, model, optimizer, criterion,
    35         optimizer.step()
    36         # Update the train loss
----> 37         train_loss = train_loss + (1 / (batch_idx + 1)) * (loss.data.item() - tr
```

```
38
39 #####
```

KeyboardInterrupt:

#### 1.1.14 (IMPLEMENTATION) Test the Model

Try out your model on the test dataset of landmark images. Use the code cell below to calculate and print the test loss and accuracy. Ensure that your test accuracy is greater than 60%.

```
In [12]: ### Do NOT modify the code below this line. ###

        # load the model that got the best validation accuracy
        model_transfer.load_state_dict(torch.load('model_transfer.pt'))
        test(loaders_transfer, model_transfer, criterion_transfer, use_cuda)
```

Test Loss: 0.923124

Test Accuracy: 75% (947/1250)

---

#### ## Step 3: Write Your Landmark Prediction Algorithm

Great job creating your CNN models! Now that you have put in all the hard work of creating accurate classifiers, let's define some functions to make it easy for others to use your classifiers.

#### 1.1.15 (IMPLEMENTATION) Write Your Algorithm, Part 1

Implement the function `predict_landmarks`, which accepts a file path to an image and an integer `k`, and then predicts the **top k most likely landmarks**. You are **required** to use your transfer learned CNN from Step 2 to predict the landmarks.

An example of the expected behavior of `predict_landmarks`:

```
>>> predicted_landmarks = predict_landmarks('example_image.jpg', 3)
>>> print(predicted_landmarks)
['Golden Gate Bridge', 'Brooklyn Bridge', 'Sydney Harbour Bridge']
```

```
In [13]: import cv2
        from PIL import Image

        ## the class names can be accessed at the `classes` attribute
        ## of your dataset object (e.g., `train_dataset.classes`)
        model_transfer.load_state_dict(torch.load('model_transfer.pt'))

        def predict_landmarks(img_path, k):
```

```

## TODO: return the names of the top k landmarks predicted by the transfer learned
loader = transforms.Compose([transforms.Resize(255),
                             transforms.CenterCrop(224),
                             transforms.ToTensor(),
                             transforms.Normalize((0.485, 0.456, 0.406),
                                                  (0.229, 0.224, 0.225))
                             ])

image = Image.open(img_path)
image = loader(image).float()
image.unsqueeze_(0)

if use_cuda:
    image = image.cuda()

model_transfer.eval()
pred = model_transfer(image)
_, top_class = pred.topk(k, dim=1)
#print(top_class) # Was testing
#indices = top_class.tolist() # Not needed at end
#print(indices) # More testing
classes = train_data.classes
#print(classes[1]) # Another 1
predictions = []
for idx in top_class[0]:
    #print(classes[idx]) # Last 1
    predictions.append(classes[idx])

return predictions

# test on a sample image
predict_landmarks('images/test/09.Golden_Gate_Bridge/190f3bae17c32c37.jpg', 5)

```

```


Out[13]: ['09.Golden_Gate_Bridge',
          '38.Forth_Bridge',
          '30.Brooklyn_Bridge',
          '28.Sydney_Harbour_Bridge',
          '06.Niagara_Falls']

```

### 1.1.16 (IMPLEMENTATION) Write Your Algorithm, Part 2

In the code cell below, implement the function `suggest_locations`, which accepts a file path to an image as input, and then displays the image and the **top 3 most likely landmarks** as predicted by `predict_landmarks`.


Some sample output for `suggest_locations` is provided below, but feel free to design your own user experience!



Is this picture of the  
Golden Gate Bridge, Brooklyn Bridge, or Sydney Harbour Bridge?

```
In [14]: import matplotlib.pyplot as plt
def suggest_locations(img_path):
    # get landmark predictions
    predicted_landmarks = predict_landmarks(img_path, 3)
    ## TODO: display image and display landmark predictions
    img = Image.open(img_path)
    plt.imshow(img)
    plt.show()
    print('The Top 3 Predicted Landmarks for this foto are: ')
    print(predicted_landmarks[0])
    print(predicted_landmarks[1])
    print(predicted_landmarks[2])

# test on a sample image
suggest_locations('images/test/09.Golden_Gate_Bridge/190f3bae17c32c37.jpg')
```



The Top 3 Predicted Landmarks for this foto are:

09.Golden\_Gate\_Bridge

38.Forth\_Bridge

30.Brooklyn\_Bridge

### 1.1.17 (IMPLEMENTATION) Test Your Algorithm

Test your algorithm by running the `suggest_locations` function on at least four images on your computer. Feel free to use any images you like.

**Question 4:** Is the output better than you expected :) ? Or worse :( ? Provide at least three possible points of improvement for your algorithm.

**Answer:**


- Well the output is as expected. The photos used from my computer were from Spain. Where no landmarks were used for training the model, so it tried to match them to similar ones in its classes.

It classified well the Golden Gate, from Google Images, but failed with the Sagrada Familia.

- Improv 1: More landmarks to train, for example from Spain, :)
- Improv 2: I was reading some discussion about lr scheduler, so I would definitely have used it to improve performance if I hadn't come up with my current attempt. Luckily wasn't necessary, but next time I will use it.
- Improv 3: Unfreeze the last layers of the pre-entrained network could help learning more advanced features from the data

- Improv 4: Different or more types or augmentation, also duplicating the images so it had more impact, 1 normal and 1 augmented


```
In [15]: ## TODO: Execute the `suggest_locations` function on
        ## at least 4 images on your computer.
        ## Feel free to use as many code cells as needed.
        suggest_locations('my_samples/055.jpg') # Ourense's Thermal Pools
```



The Top 3 Predicted Landmarks for this foto are:

```
03.Dead_Sea
02.Ljubljana_Castle
48.Whitby_Abbey
```

```
In [16]: suggest_locations('my_samples/57.jpg') # Sagrada Familia Barcelona
```




The Top 3 Predicted Landmarks for this foto are:

19.Vienna\_City\_Hall

16.Eiffel\_Tower

28.Sydney\_Harbour\_Bridge

In [17]: `suggest_locations('my_samples/94.jpg')` # *Golden Gate*






The Top 3 Predicted Landmarks for this foto are:

09.Golden\_Gate\_Bridge

30.Brooklyn\_Bridge

38.Forth\_Bridge

```
In [18]: suggest_locations('my_samples/2905.jpg') # Catedral of Santiago de Compostela
```



The Top 3 Predicted Landmarks for this foto are:

10.Edinburgh\_Castle

21.Taj\_Mahal

19.Vienna\_City\_Hall

## 2 *Stand Out Suggestion 3*

Additional use cases: I read online that VGG is being used from numerous multilabel classifications of all kind from brain diseases to clothing material, and even to facial attributes. So it may be useful for a huge number of possible use cases, as can be seen in these examples. One could be, helping police agents to identify criminal photos and classify their location based on photos taken from their social media site. Similar to what we have done here with landmarks but in a higher quality.